

Property Testing Pandas with Bulwark

Zax Rosenberg, CFA
Senior Data Scientist @ SPINS
github.com/ZaxR
www.zaxrosenberg.com





Agenda

- About Me
- About Property Testing
- About Bulwark
- Demo
- Contributing



About Me

- Senior Data Scientist @ SPINS
- Director of ChiPy's Mentorship Program
- Co-host of ChiPy's Project Night
- Blogger (zaxrosenberg.com/blog)
- Speaker (github.com/ZaxR/talks)



About Me

- NOT a testing expert



What is Property Testing?

Checking that some object has certain properties.
For example:

```
some_list = [1, 2, 3, 4]
```

One property of `some_list` is that its values are in a range of 1-4. Another is that it's mutable.



Why is it valuable?

- Testing reduces bugs/lowers development cost
- We don't always have the exact data up front
- Property tests can be relatively fast to run
- Easy to include domain knowledge



Introducing Bulwark



Introducing Bulwark

- Bulwark is an open-source library that lets you easily property test pandas dataframes.
- It's designed to make it easier for data analysts and data scientists to test our own code.



Bulwark's Design


- Property tests are available as functions (“**checks**”) and **decorators**.
- Each check:
 - Takes a `pd.DataFrame` and optional additional arguments,
 - Makes an assertion about the `pd.DataFrame`, and
 - Returns the original, unaltered `pd.DataFrame`.
- A failed checks raises an **`AssertionError`**, printing an informative message.
- Each **check** has an auto-magically-generated associated **decorator**, allowing you to make your assertions *outside* the actual logic of your code. This is a core benefit of Bulwark.



Quickstart - Input

```
In [1]: import bulwark.decorators as dc ← import
...: import numpy as np                convention
...: import pandas as pd
...:
...:
...: @dc.HasNoNans() ← a check
...: def add_five(df):                 decorator
...:     return df + 5
...:
...:
...:
...: df = pd.DataFrame({"a": [1, 2, 3], "b": [4, 5, np.nan]})
...: add_five(df)
```

uh oh...





Quickstart - Result

```
~/Projects/bulwark/bulwark/checks.py in has_no_nans(df, columns)
    99
   100     """
--> 101     return has_no_x(df, values=[np.nan], columns=columns)
   102
   103

~/Projects/bulwark/bulwark/checks.py in has_no_x(df, values, columns)
    77
    78     try:
--> 79         assert not df[columns].isin(values).values.any()
    80     except AssertionError as e:
    81         missing = df[columns].isin(values)

AssertionError: (2, 'b') ← row index 2, column 'b' fails
```



What if I have multiple checks?

```
In [2]: import bulwark.checks as ck
...: import bulwark.decorators as dc
...: import numpy as np
...: import pandas as pd
...:
...:
...: @dc.MultiCheck(checks={ck.has_no_nans: {"columns": ['b']},
...:                               ck.is_shape: {"shape": (3, 1)}})
...: def subtract_five(df):
...:     return df - 5
...:
...:
...: df = pd.DataFrame({"a": [1, 2, 3], "b": [4, 5, np.nan]})
...: subtract_five(df)
```

pass a dict of
check: param pairs

What if I have multiple checks?

```
~/Projects/bulwark/bulwark/checks.py in multi_check(df, checks, warn)
    434         return df
    435     elif error_msgs:
--> 436         raise AssertionError("\n".join(str(i) for i in error_msgs))
    437
    438     return df
```

AssertionError: (2, 'b')

Expected shape: (3, 1)

Actual shape: (3, 2)




collects all the errors



What if I don't want to raise errors?

```
In [3]: import bulwark.checks as ck
...: import bulwark.decorators as dc
...: import numpy as np
...: import pandas as pd
...:
...:
...: @dc.MultiCheck(checks={ck.has_no_nans: {"columns": ['b']},
...:                               ck.is_shape: {"shape": (3, 1)}},
...:               warn=True)
...: def subtract_five(df):
...:     return df - 5
...:
...:
...: df = pd.DataFrame({"a": [1, 2, 3], "b": [4, 5, np.nan]})
...: subtract_five(df)
```



prints instead of raises



What about when I go to production?

```
In [4]: import bulwark.decorators as dc
...: import numpy as np
...: import pandas as pd
...:
...:
...: @dc.IsShape((3, 2), enabled=False)
...: def subtract_five(df):
...:     return df - 5
...:
...:
...: df = pd.DataFrame({"a": [1, 2, 3], "b": [4, 5, np.nan]})
...: subtract_five(df)
```

← turn off this check

← doesn't raise an error

Pro tip: set a centralized config variable that toggles all decorators' statuses.



Demo Time!



Where should I use Bulwark?

- On ETL pipeline functions, especially E & L
 - Help enhance your understanding of the data upfront, even if you don't do full EDA
 - Integration test by checking output
- In unit tests



How should I use Bulwark?

- Favor the decorator version within core code
 - Lets you disable/switch to warnings
 - Separates checks from code/business logic
- Use check version in unit tests



When should I use Bulwark?

- During development
- Maybe at run time.
 - If it's acceptable for runs to fail
 - If the state should never be reached, and you can't handle the error.



Who's using Bulwark?

- 6,123 total downloads
- 256 downloads/month (excluding mirrors)
- Folks in > 45 countries



Contributing is easy, too!

- Very friendly to folks new to open source!
Adding a new check is as easy as writing a single function.
- Full instructions available at:
<https://bulwark.readthedocs.io/en/latest/contributing.html>



Find out more

- PyPI: <https://pypi.org/project/bulwark/>
- Read the Docs: <https://bulwark.readthedocs.io>
- GitHub: <https://github.com/ZaxR/bulwark>